Sample Document for RAG Loader Testing

Introduction

Retrieval-Augmented Generation (RAG) is a technique that combines information retrieval with text generation. Instead of relying only on information stored in a language model, a RAG system can retrieve relevant passages from an external knowledge source.

How a Basic RAG Pipeline Works

A basic RAG pipeline has several stages. First, a document loader reads source files and extracts their text. Next, a chunker divides the text into smaller sections. An embedding model converts each chunk into a numerical vector. The vectors are stored in a vector database, where a retriever can search for relevant information when a user asks a question.

Document Loading

A document loader is responsible for reading a supported file format and extracting its content. For a PDF, a loader may extract text page by page. Keeping page numbers alongside the extracted text makes it easier to identify where retrieved information came from.

Chunking

Chunking splits long text into smaller pieces. Smaller chunks can make retrieval more focused because the system can return a relevant passage rather than an entire document. A chunking strategy may use a fixed character length, sentence boundaries, or token counts. Some strategies also use overlap so that context near a chunk boundary is not lost.

A Simple Test Fact

The fictional project library in this sample contains 12 documents. This fact is included only to provide a clear sentence that can be found when testing text extraction.